
Image Processing

Licence Plate Detector



1 Introduction (0.5 pt)

We start our program in the `CaptureFrame_Process` with following parameters. The `file_path`, the file path to the video, the `sample_frequency`, how many seconds we have between 2 frames we are going to analyse and the `save_path`, where we're storing the file in the final step. Our program starts by reading in the video, if we cannot find the video we print "No video found!" and terminate the program.

We extract some usefull information from the video (e.g. `total_frame` and `fps_video`) which we will use later on. We loop over all the frames and analyse every frame where the frame number modulus `frames_to_skip` is 0. `Frames_to_skip` is calculated using the `sample_frequency`. From there we first localise the number plate, than optimize the found number plate (e.g. rotation). Next we split the optimized number plate into different images such that every image contains a number/letter/noise. We than filter out the noise and recognize every symbol by comparing it to each symbol in a dataset.

In *figure 1* is the original image from the training video. This we used for training and for this document will be used to work on.



Figure 1: Original Image

2 License plate localization method (4.0 pt)

2.1 License plate localization data description (0.5 pt)

For the localization we did not need to look up a lot. We looked up the HSI values for yellow and in which range these values are. We also looked up how you can combine an image and a mask such that you only keep the part of the image that is also in the maks. The rest was mainly taken from what we learned from the lectures. We found our information for the localization on stack overflow and in the lecture slides. From there we combined the sources of information to create the `plate_detection` in our localization.

24 From the data we were given by the course we chose to use "trainingsvideo.avi" as our test
25 video, but we skip the first 36 frames as these frames are already in our training data. Those 36 frames
26 are exactly the "Video10_2.avi" in the TrainingSet provided by the course. After applying
27 the function we made we printed out the before and after image and the data. This way we could
28 manually label if they were a pass or fail. We could check what the problem was directly and note it
29 down as well.

30 For the optimization we used "Video10_2.avi" which was 36 frames with the same number plate.
31 As soon as we were satisfied with the result of that number plate we performed a test on 29 number
32 plates from different categories which were not used for training so we skipped the first 36 frames of
33 the "trainingsvideo.avi" video. We counted a test as a pass if and only if at least in 1 frame the
34 full licence plate and potentially some noise was returned by the method.

35 2.2 License plate localization system (1.5 pt)

36 The localization method takes in a frame and starts working on that until we are only left with the
37 licence plate and maybe a bit of noise. The area of the noise « area of the licence plate. We start by
38 creating a mask for yellow values, because our first focus was Dutch number plates who are always
39 yellow for regular road cars. This mask you can see in *figure 2*. Before clearing out the holes in the
40 mask we first find all corners and save their coordinates for later. Then we dilate 10 times over the
41 mask to remove the holes in our mask where the letters should be and to make sure we get the full
42 licenceplate. We still have the coordinates of the corners and we use those later to cut out the perfect
43 licenceplate, but before we can use those we first need to rotate the plate. Here we return the plate
44 with noise (because we gonna use that noise for rotation) and the 4 corner values to later cut out the
45 licenceplate after rotating. The rotation is done by first finding a line that minimizes the distance
46 from every point in the mask to the line. We calculated this using the least-squares-method. After the
47 line is found it is represented as $mx+b=y$. If we take the atan of m we can find the angle at which the
48 plate stands.



Figure 2: The mask before and after dilation

49 2.3 Evaluation metric for license plate localization (1.5 pt)

50 For the evaluation we used the test video that is described above. Checking if a plate was found was
51 quite easy because we only kept the licence plate part. By simply displaying the found licence plate
52 we could just say pass or fail by looking at our definition of a passing licence plate. This gives us
53 also extra insights, if the plate was a fail we could note it down. We can use the gained information to
54 improve our solution. We counted a test as a pass if and only if at least in 1 frame at least the full
55 licence plate and potentially some noise was returned by the method. We know that "some noise" is a
56 bit vague, but we counted it as too much noise, if it is dominant over the found plate. A small border
57 of noise around it is allowed, which was even very handy to find the rotation later.

58 We end up with a couple of numbers. The number of plates, the number of correct found and the
59 number of wrongly found plates. To calculate the precision we do $precision = \frac{\#correctPlates}{totalplates}$. It is a
60 simple formula, but sufficient enough for this problem.

61 2.4 Analysis of the license plate localization results (0.5 pt)

62 We get the following results:

63 We find 28/29 licence plates in Cat I

64 We find 10/10 licence plates in Cat II

65 We find 10/10 so 5/5 licence plates in Cat III

66 We find 0/10 licence plate in Cat IV

67



Figure 3: An example of a correct found licence plate



Figure 4: An example of an incorrect found licence plate

68 3 Character recognition method (4.0 pt)

69 3.1 Recognition data description (0.5 pt)

70 For the training of the segment-and-recognize part we used the first number plate of the video
71 provided by the course. The model was validated using a small dataset we made by searching images
72 of number plates on google. This data was used for both the segmenting and recognizing of the
73 number plates/symbols. The actual recognition is done by comparing each symbol to each character
74 in the dataset we got from the course. However we found out this was not sufficient enough so we
75 made our own letters/numbers for most of given letter/numbers. This was made by our method of
76 cutting out letters/numbers of our licence plate and simply saving that result in a bmp file. We used
77 both the course and our dataset. We were happy with our method when we got an output in ascii that
78 exactly matched the symbols in our number plate. We checked this manually by printing the output
79 to the console and showing the input using imshow(). We did not count any partial correct output as
80 the eventual goal is to recognize complete number plates.

81 3.2 Recognition system (2.0 pt)

82 3.2.1 Character localization system (1.0 pt)

83 The Localization of our character recognition process has 5 concrete steps. We first start with rotating
84 the image over the angle given to us by the Localization method. After this step we have an image of
85 a number plate with next-to-no noise surrounding the number plate. The numberplate itself is straight.
86 We first convert the image to a greyscale image. This will help us out later when we convert it to a
87 binary image. The conversion to a binary image is done by first calculating the ISODATA threshold.
88 We then set everything that's below this threshold to 255 and everything above to 0. This way the dark
89 parts of the image, which are the symbols in the number plate, will be white and the background will
90 be black. Afterwards we dilate and erode the image to both get rid of noise and fill in the gaps in the
91 symbols. The final actual localization step of our characters is done by iterating over the columns of
92 our image. We determine a certain threshold. For each column we check if there are less white value
93 than this threshold, if this is the case we can assume this column is the end of a character and we
94 split our image on that column.



Figure 5: input vs output

95 3.2.2 Character recognition system (1.0 pt)

96 Our input is an image of variable size which can contain either a symbol, or a line. We first do some
97 preprocessing on the input image. This contains normalizing the size of the image and cropping the
98 image so we don't have any completely black rows or columns. Next we start a for-loop through all
99 the images in our character dataset. For each of these images some preprocessing is needed: they are
100 not complete binary images so we convert them to binary. We also resize them to the same size as
101 our input image. Finally we can start comparing them, we do this by taking the distance between the
102 input image and each of the dataset images. We then take the image with the lowest distance to be
103 our symbol. If this distance isn't low enough we assume the input image to be either noise or a line.

104 3.3 Evaluation metric for recognition (1.0 pt)

105 We made the method by looking how it worked on the first number plate in the video. After we
106 were happy with that result we switched to the dataset we got from google. We used the metric:
107 $\frac{\#fullplatesRecognized}{\#allplates}$. Like our metric from the localization this is a simple but sufficient metric.
108 We also use this metric because we only have a correct plate when we recognize each symbol. We
109 checked the results manually, by comparing the outputs the method gave us and the number plates
110 themselves. When we were finally happy with the result we moved on to optimization. From the
111 dataset found on google we recognized the full 10 out of 10 number plates.

112 3.4 Analysis of the recognition results (0.5 pt)



Figure 6: input vs output

113 One case where our method really struggles is if there are a lot of shadows. This can best be seen in the
114 number plate shown above. We use the ISODATA threshold to separate the foreground: symbols
115 from the background. If there are a lot of shadows the plate gets too dark to separate the symbols
116 from background as the color difference between the 2 is not there. This causes the binary image to

117 have a lot of noise at the top of our image. Subsequently this causes the splitting of the symbols to
118 not work properly as it cannot find an x-value where the symbols should be split. The examples given
119 above in figure 5 and 6 are a perfect example of this problem. We could improve this method by
120 somehow making the method invariant to illumination. This by for example using HSI values instead
121 of RGB values. Due to time limitations we cannot complete this.

122 **4 Analysis of system (1.5 pt)**

123 **4.1 Time analysis (0.1 pt)**

124 After going over the system with timers we found that the system is the slowest in recognition part.
125 This due to the comparison of 28 letters and 18 numbers. This takes some time. That is why we only
126 analyze with a frame rate 1/2 in stead of the default frame rate of 1/12.

127 **4.2 Successes and failures (1.0 pt)**

128 Simple plates that do not have much in common with the car (so no white or black cars) preform the
129 best. Due to the ISODATA we are very vulnerable for these kinds of cars. We are also vulnerable for
130 shadows and see that this gives us problems in the recognition. This is due to the choose to do mostly
131 work on the color image and not for example edge detected images.

132 **4.3 Future improvements (0.4 pt)**

133 An improvement would be to use edge detection instead of a color detection to Localize the plates.
134 This way we would also be able to find international plates as well. Another aspect we need to
135 improve is speed. A the moment of writing this we take about 1 - 1.5 seconds for each frame. The
136 way we could do this is by removing a lot of double for loops in the system. The localizing of the
137 symbols could also be improved by using blob detection instead of the current method.